

Classification with Machine Learning

Kevin Wang

Colaboratory notebook:

<https://colab.research.google.com/drive/1Cein0r-J9N2vX1xh24cRLEHgBkJx3p7w>

You should know the following things

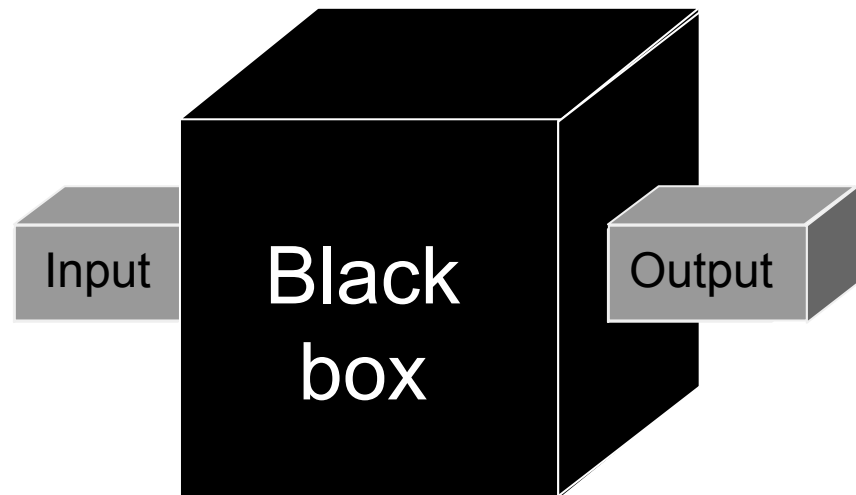
- Basic matrix operations
- Derivatives
 - If you don't know – just think of them as slope of function at any given point
- Basic programming skills
 - We use Python

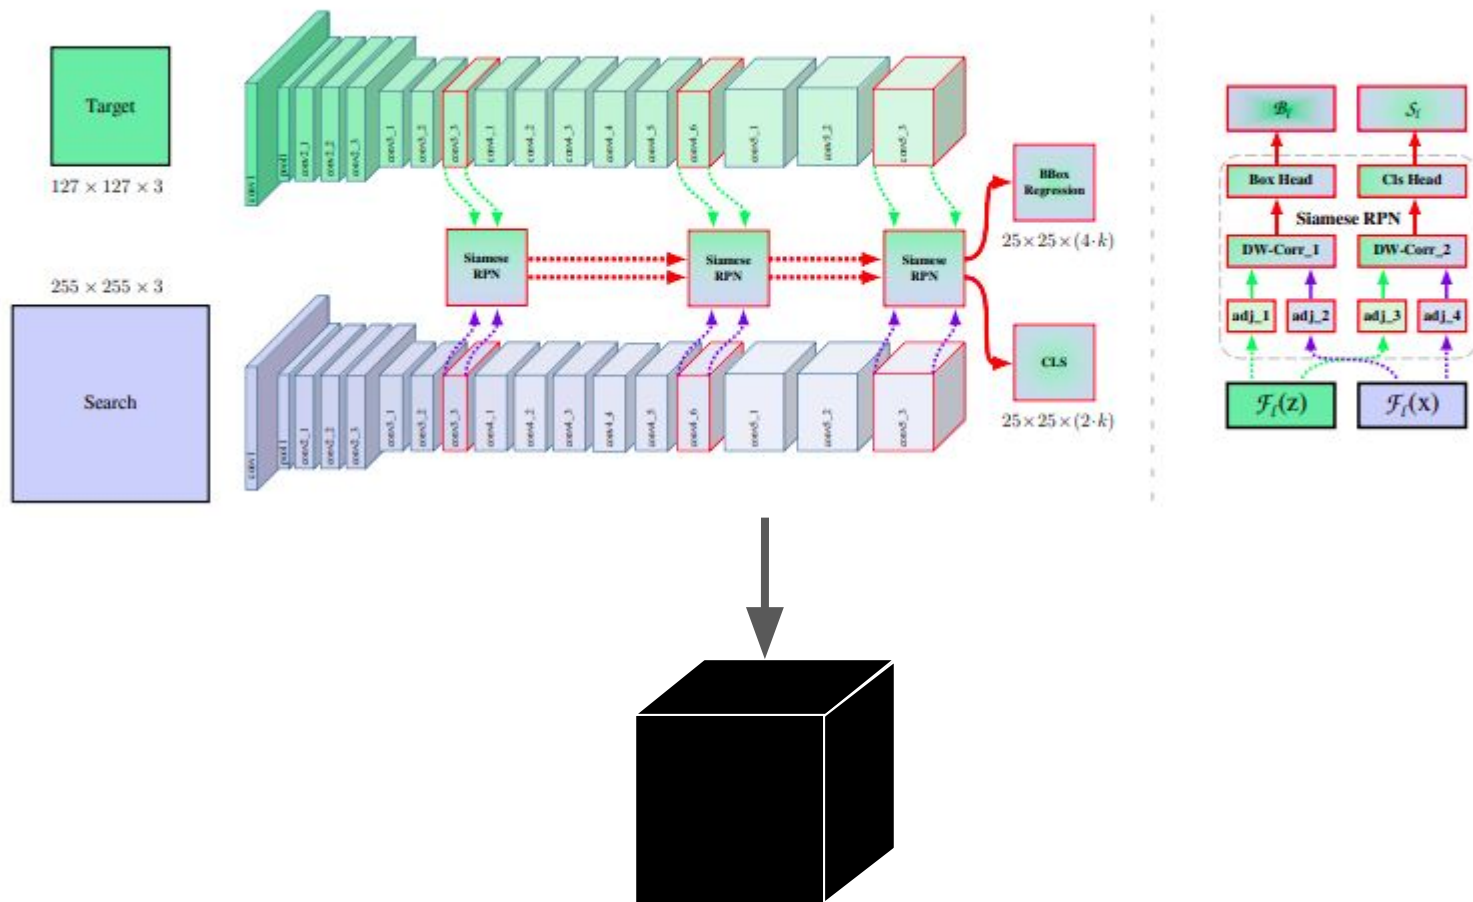
What is machine learning?

- “Field of study that gives computers the ability to learn without being explicitly programmed”
 - - Andrew Samuel, ML pioneer in the 50s
- You create a model that is trained to perform some task
 - Model: an algorithm that can be adjusted by training
 - Training: the stage where the model learns how to perform a task
- Almost all artificial intelligence research today is in the field of machine learning
 - Other AI methods exist, but are not as powerful / promising

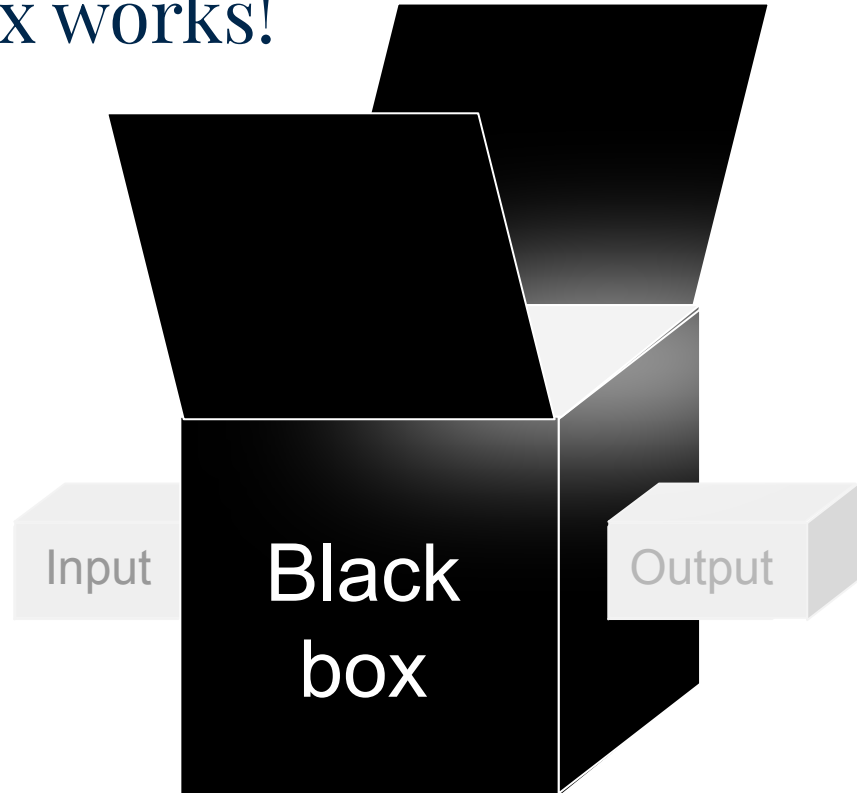
Another way to think of it: black box model

- You can see the input and the output
- A black box takes in the input to produce the output
 - You don't need to configure the inside of the black box to make use of it
- You define the model in your black box using machine learning
 - The power of ML is that we don't need to configure the black box manually!





The purpose of these tutorials is to learn how the inside of the black box works!



Classification

- Input: data
 - Image of UMich logo
 - Information about a tumor
 - Post from Piazza
- Output: what class is the input?
 - Is it a UMich logo or a State logo?
 - Is tumor cancerous or benign?
 - Is this post about an exam or a project?

Classes of input data are their labels

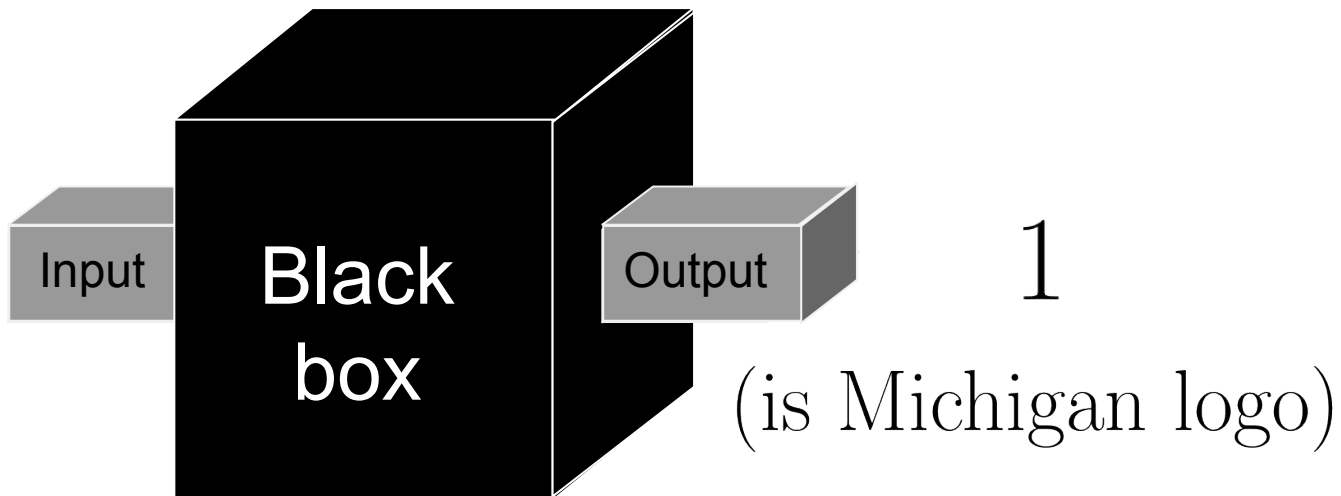
Class = 1 (is a Michigan logo)



Class = 0 (is not a Michigan logo)



We'd like the model to do this:

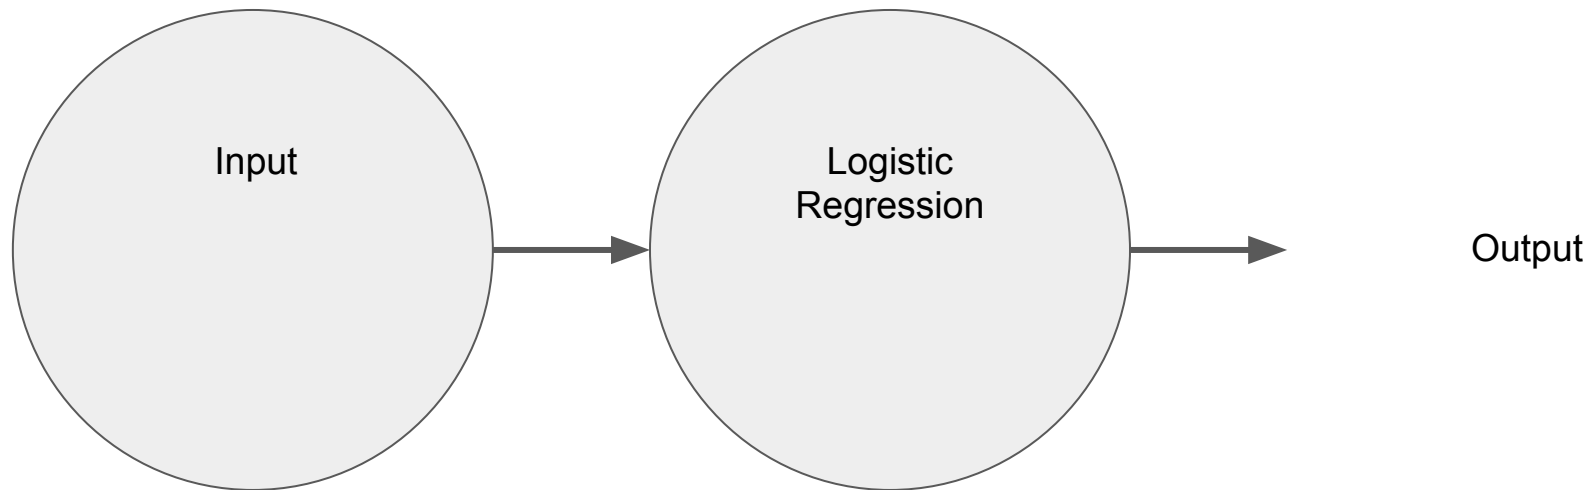


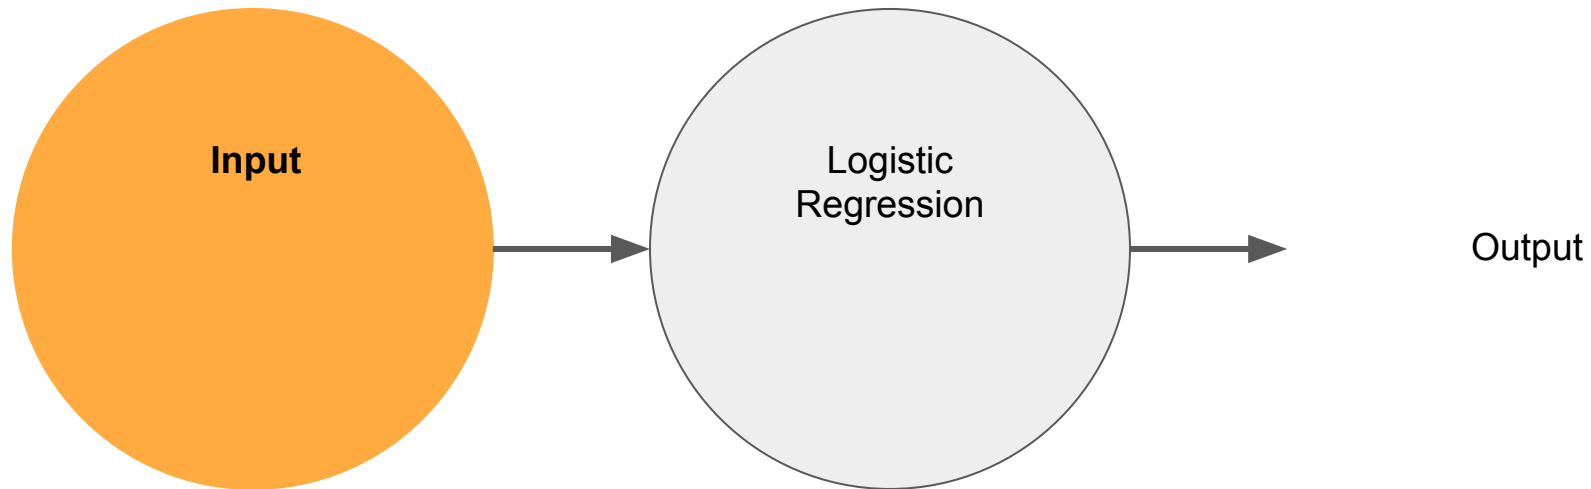
How?

- For today, we're using a model that will perform logistic regression
 - “logistic” indicates the function that makes the decision...more on that later

The flow of logistic regression:

- Training step: requires labeled training data
 - Format the input
 - Break it down into features
 - Several times:
 - Run the model, get your output
 - Based on your output, revise the model
- Test step: input data, run model, and get output
 - Note there is no labeling in this step
- This is similar to how other ML models work





The training dataset contains many instances / examples of data

- Use lowercase x for a single instance
- Denote which instance using superscript with parentheses
 - i is the standard placeholder variable
- We say there are m total
 - Not a universal convention!
- Case and superscript are important!
 - Leaving out superscript usually means arbitrary instance

This is the third instance / example in the training dataset

$$x^{(3)}$$

To represent an arbitrary instance, use

$$x^{(i)}$$

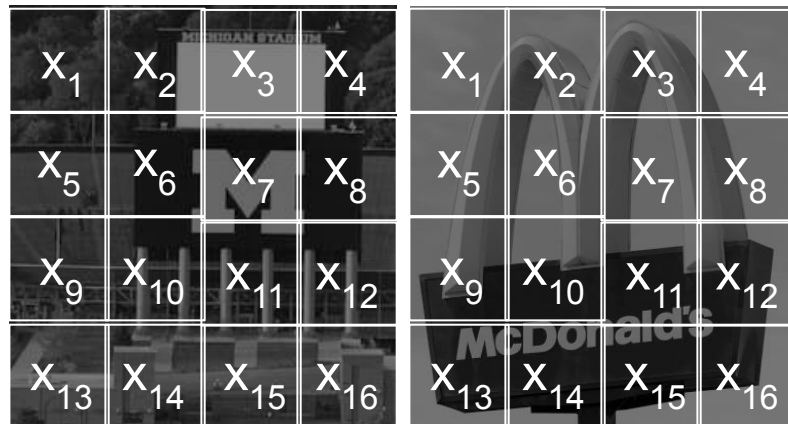
or

$$x$$

Defining features

- Features might be:
 - pixel values (images)
 - size, age, location, growth rate (tumor)
 - words in post (Piazza)
- Features are denoted by subscript:
 - x_1, x_2, x_3, \dots
 - Indices start at 1
- Every instance has the same number of features
 - This number is n

Breaking down two images into
several features based on pixel
value



(pixels not drawn to scale)

Instance with its features is represented by a vector

In general, this is how a single instance
is represented

$$x^{(i)} = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ \dots \\ x_n \end{bmatrix}$$

- How might we represent the dataset as a whole?

The dataset as a whole is represented by a matrix

- Name is X
 - vector = lowercase, Matrix = Uppercase
- Instances are stored in a row
- One column is one instance (with its features)
 - Rows are less useful for us
- It can often help to think of X as a row vector of column vectors
 - Multiplication

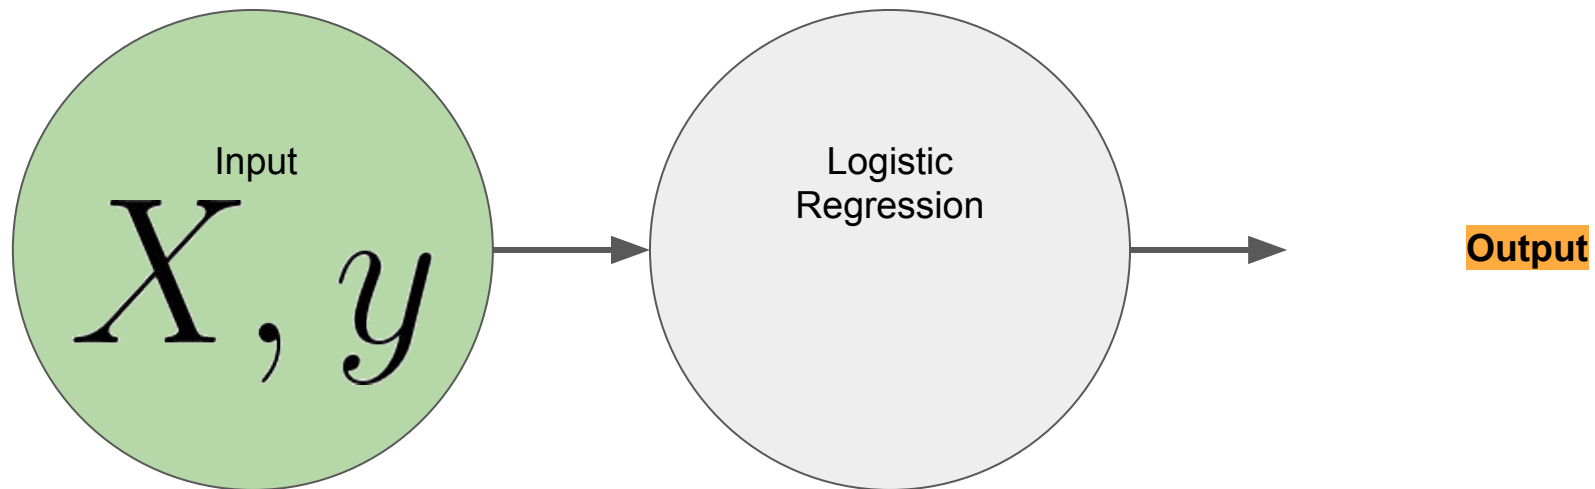
$$X = \begin{bmatrix} | & | & | & \dots & | \\ x^{(1)} & x^{(2)} & x^{(3)} & & x^{(m)} \\ | & | & | & & | \end{bmatrix}$$

Don't forget the labels:

- Also stored as a vector
 - A row vector this time
- Superscript in parentheses denotes corresponding example $x^{(i)}$ from the dataset

Note superscripts to denote specific examples

$$y = \begin{bmatrix} y^{(1)} & y^{(2)} & y^{(3)} & \dots & y^{(m)} \end{bmatrix}$$



The output is a vector of predicted probabilities

- Output is a vector of predicted classes for corresponding instances in X
 - We call this vector " y_hat "
 - Also a row vector
- At test time, you'd round the predicted values
 - Training uses unrounded values

The i th value in y_hat equals the predicted probability (given the i th training example) that

$$y^{(i)} = 1$$

$$\hat{y}^{(i)} = \text{predicted } P(y^{(i)} = 1 | x^{(i)})$$

y_hat is used to denote the entire vector

$$\hat{y} = \begin{bmatrix} \hat{y}^{(1)} & \hat{y}^{(2)} & \hat{y}^{(3)} & \dots & \hat{y}^{(m)} \end{bmatrix}$$

And now a brief digression...

...to matrix dimensions

Pop quiz:

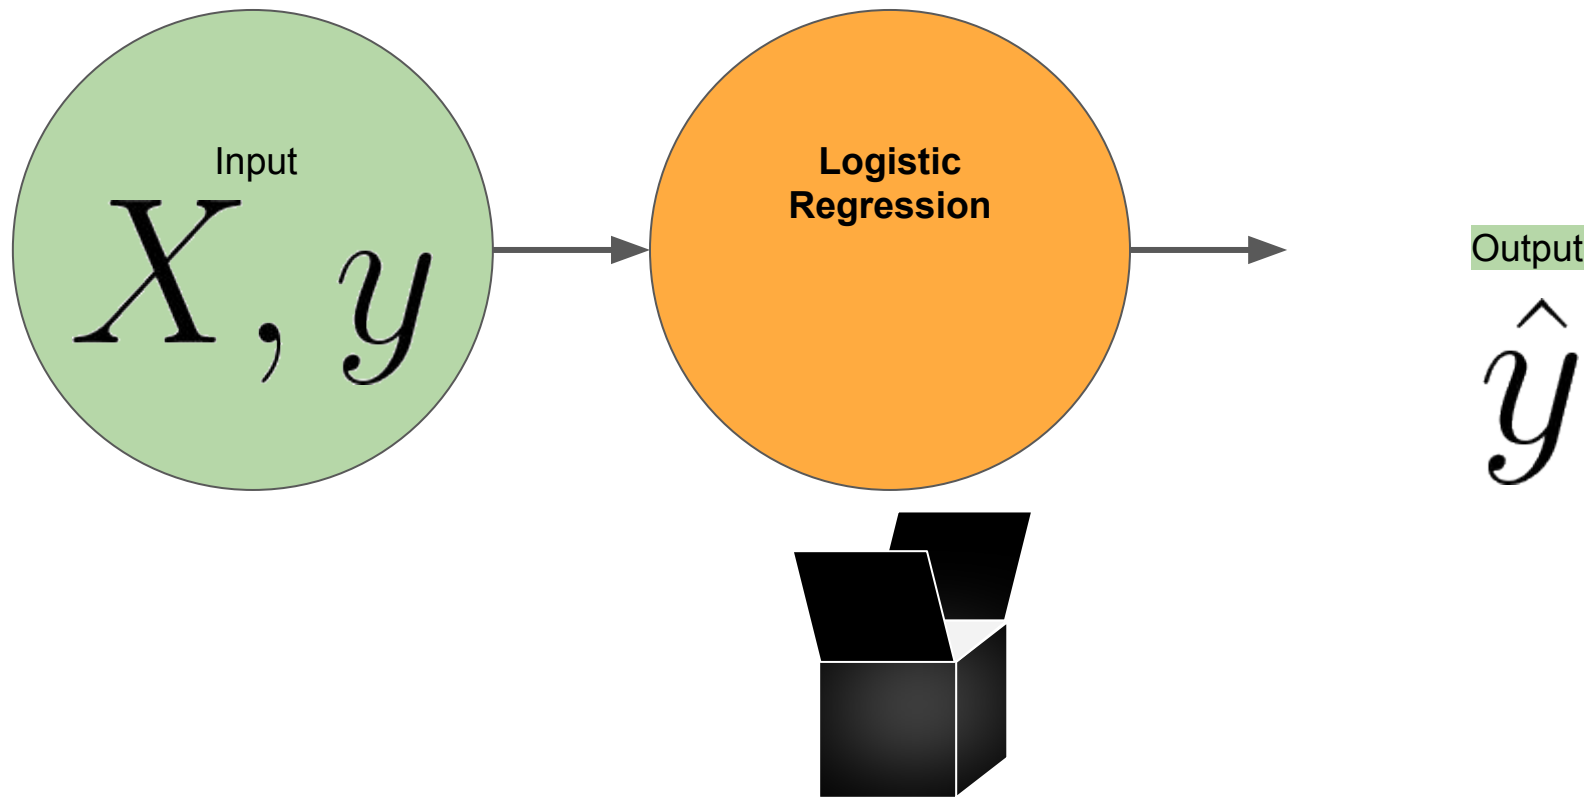
- What are the dimensions of y ?
 -
- What are the dimensions of y_{hat} ?
 -
- What are the dimensions of a single training example $x^{(i)}$?
 -
- What are the dimensions of X ?
 -
- Why should every training instance have the same number of features?
 -
- Let m = number of instances and n = number of features

Pop quiz:

- What are the dimensions of y ?
 - $(1, m)$
- What are the dimensions of y_{hat} ?
 - $(1, m)$
- What are the dimensions of a single training example $x^{(i)}$?
 - $(n, 1)$
- What are the dimensions of X ?
 - (n, m)
- Why should every training instance have the same number of features?
 - 1. it wouldn't make sense otherwise!
 - 2. keep dimensions of $x^{(i)}$ consistent for all i
- Take questions

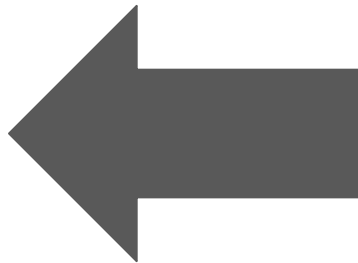
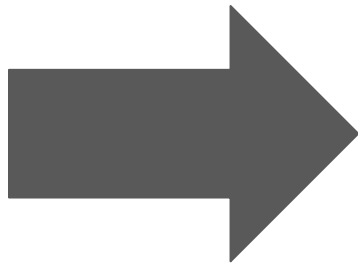
Vectors and matrices simplify our math

- And they will help with the intuition for logistic regression
 - The inner workings of the black box...
 - ...and the interesting part of today



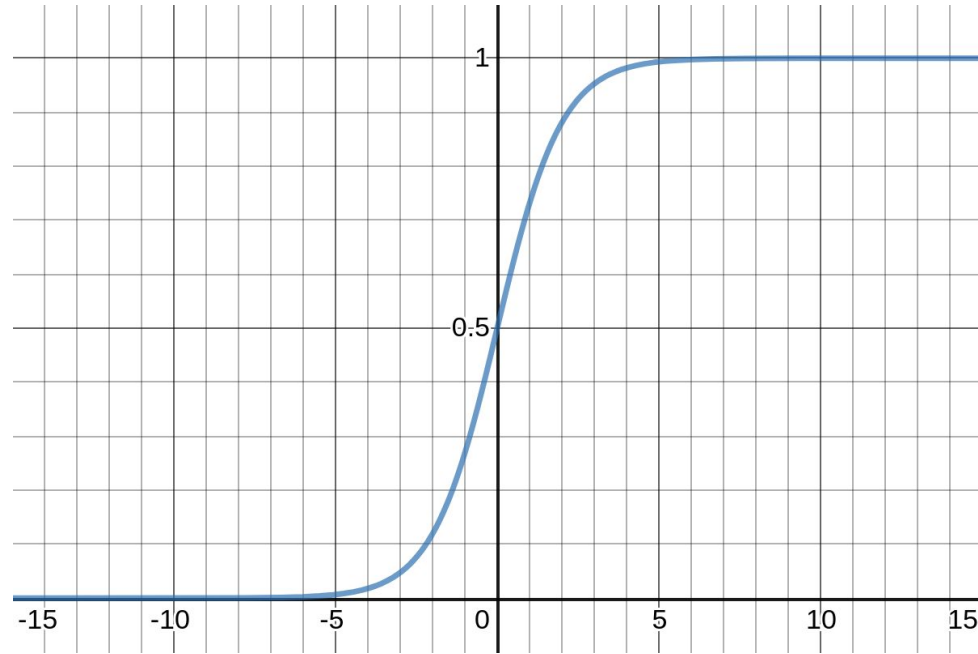
There are two steps in logistic regression

- First step: forward propagation
 - Given matrix of instances X , make predictions on all of them
- Second step: backward propagation
 - Calculate derivatives
 - Update values used in the forward propagation step



First step: forward propagation

The sigmoid / logistic function makes the predictions

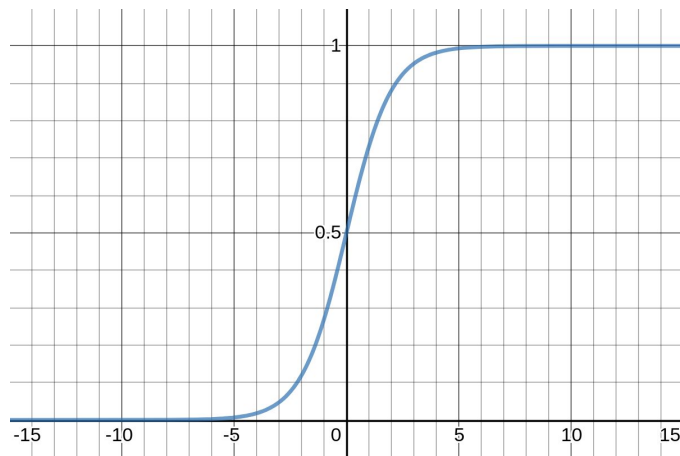


Sigmoid function: $\hat{y}^{(i)} = \sigma(x^{(i)}) = \frac{1}{1 + e^{-x^{(i)}}}$

Why is the sigmoid function a good choice?

- For all possible values x , outputs values between 0 and 1
 - This will be our predicted probability that x is class 1
- The more extreme the value of x is, the more certain the output
 - Very positive x : $\hat{y} \approx 1$
 - Very negative x : $\hat{y} \approx 0$

$$\hat{y}^{(i)} = \text{predicted } P(y^{(i)} = 1|x^{(i)})$$



$$\hat{y}^{(i)} = \sigma(x^{(i)}) = \frac{1}{1 + e^{-x^{(i)}}}$$

Why won't this work?

Why won't this work?

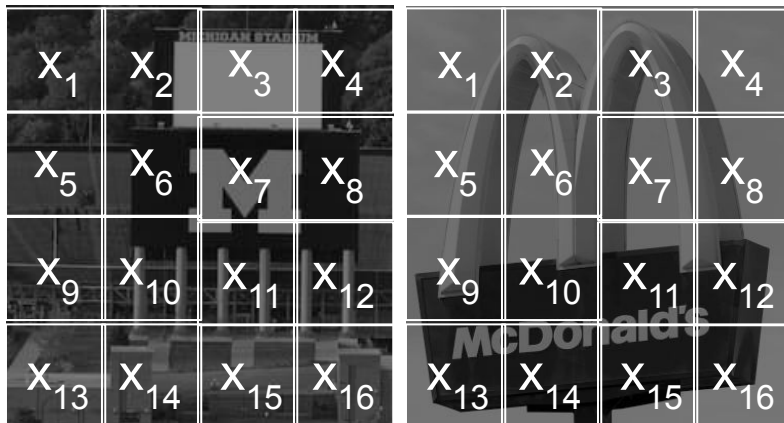
- Setting a scalar equal to a vector doesn't make sense

$$\hat{y}^{(i)} = \sigma(x^{(i)}) = \frac{1}{1 + e^{-x^{(i)}}}$$

- $\hat{y}^{(i)}$ is a scalar...
 - ...but $x^{(i)}$ is a vector

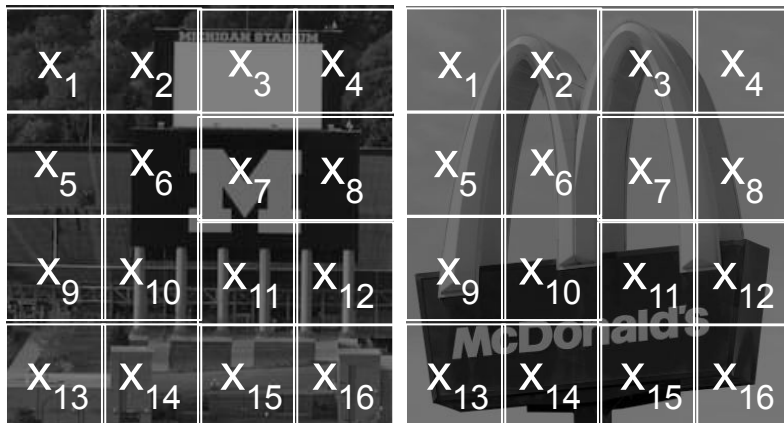
Why won't this work?

- Do all features have equally important information?
 - Are the border pixels of an image as important for classification as the center pixels?
 - Does the word “the” indicate as much about a Piazza post as the word “midterm”?



Why won't this work?

- Do all features have equally important information?
 - Are the border pixels of an image as important for classification as the center pixels?
 - Does the word “the” indicate as much about a Piazza post as the word “midterm”?
- Other way of wording it: do all features have the same weight?



Our goal is to take a weighted sum of the features

- This will be a scalar z
 - So we can properly calculate scalar $y_hat^{(i)}$
 - We'll need to calculate a $z^{(i)}$ for every $x^{(i)}$
- We'll need to represent our weights somehow

$$z^{(i)} = w_1 x_1^{(i)} + w_2 x_2^{(i)} + \dots + w_n x_n^{(i)}$$

The weight vector

- Multiply every feature by a weight
 - Weights are stored in vector w
 - Each feature has a corresponding weight

$$x_1 \rightarrow w_1 x_1$$

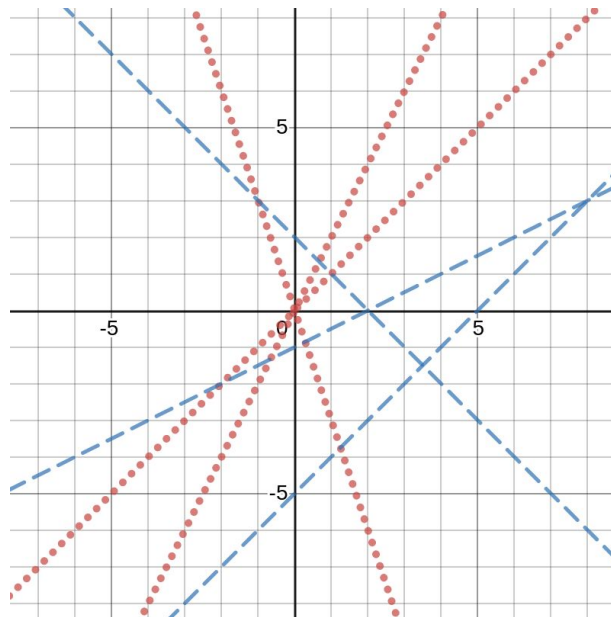
$$x_2 \rightarrow w_2 x_2$$

...

$$x_n \rightarrow w_n x_n$$

The bias

- Add a bias value to the weighted sum $z^{(i)}$
 - Bias is denoted as b
 - Think of it like y-intercept versus no y-intercept – you have more options for your model



Representing weights and bias

- We just need one bias term b
 - b is not a vector but a scalar
- What are the dimensions of w ?
 - How many weights do we need?

Representing weights and bias

- You have to have a weight value for every feature, but just add one bias value
 - So there are n weights, 1 bias
- The formula, with z being the result:
 - Note $w^T x^{(i)}$ is the same as $w_1 x_1^{(i)} + w_2 x_2^{(i)} + \dots$

$$z = w^T x + b$$

Weight vector:
multiply

$$w = \begin{bmatrix} w_1 \\ w_2 \\ w_3 \\ \dots \\ w_n \end{bmatrix}$$

Bias scalar: add

b

- We should technically be clearer here – how?

One value of z per instance, and a z vector

- The i th z value is denoted with same parenthetical superscript as $x^{(i)}$, $y^{(i)}$

Capital z is a vector of size m :

$$z^{(i)} = w^T x^{(i)} + b$$

$$z = \begin{bmatrix} z^{(1)} & z^{(2)} & z^{(3)} & \dots & z^{(m)} \end{bmatrix}$$

The beauty of matrices

- The equation for entire vector z is easy to compute using matrix multiplication
 - Python has libraries to work with matrices

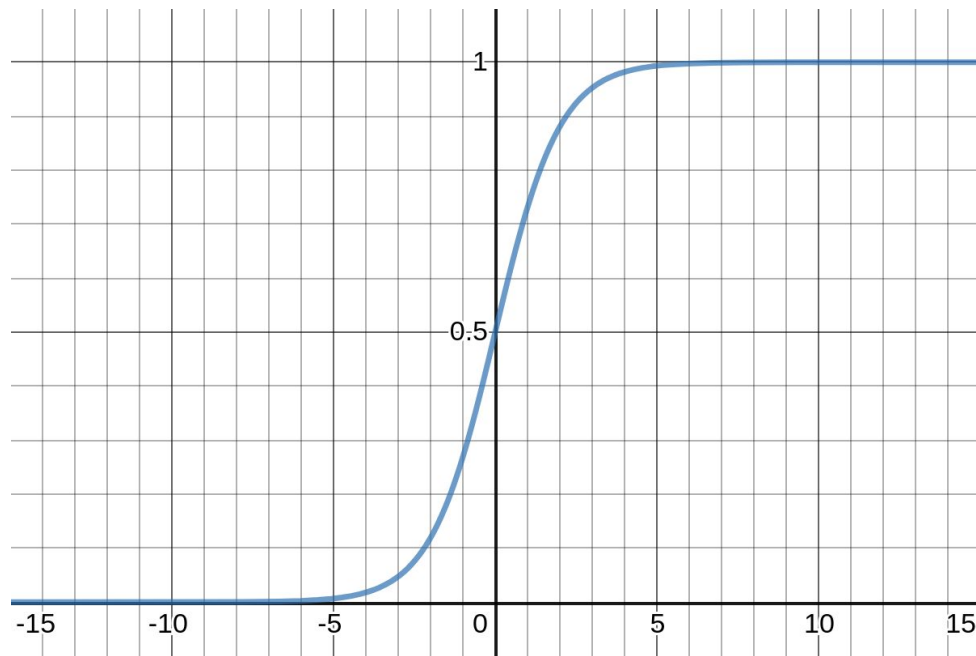
$$z = w^T X + b$$

The beauty of matrices

- The equation for entire vector z is easy to compute using matrix multiplication
 - Python has libraries to work with matrices

$$\underset{(1, m)}{z} = \underset{\substack{(n, 1)^T \\ = (1, n)}}{w^T} \underset{(n, m)}{X} + \underset{\substack{(1, 1) \\ \text{added } m \\ \text{times}}}{b}$$

The i th value of \hat{y} is computed on i th value in z



Sigmoid function: $\hat{y}^{(i)} = \sigma(z^{(i)}) = \frac{1}{1 + e^{-z^{(i)}}}$

Detail: sigmoid function is an activation function

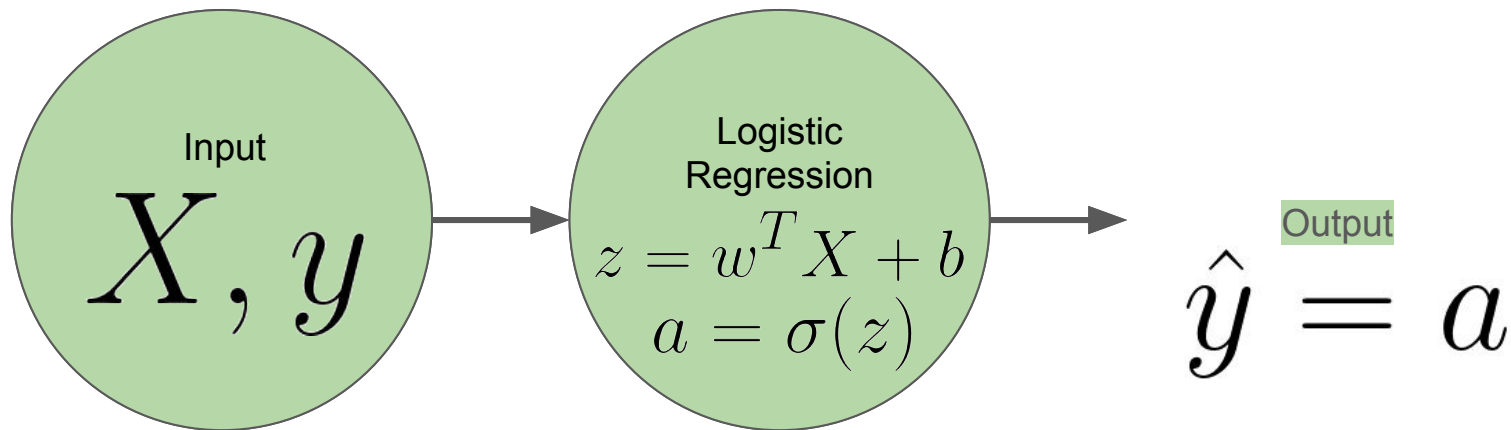
- “How much does z ‘activate’ the function?”
- y_{hat} will also be called a , and $y_{\text{hat}}^{(i)}$ will also be called $a^{(i)}$

$$a = \hat{y}$$

$$a^{(i)} = \hat{y}^{(i)} = \frac{1}{1 + e^{-z^{(i)}}}$$

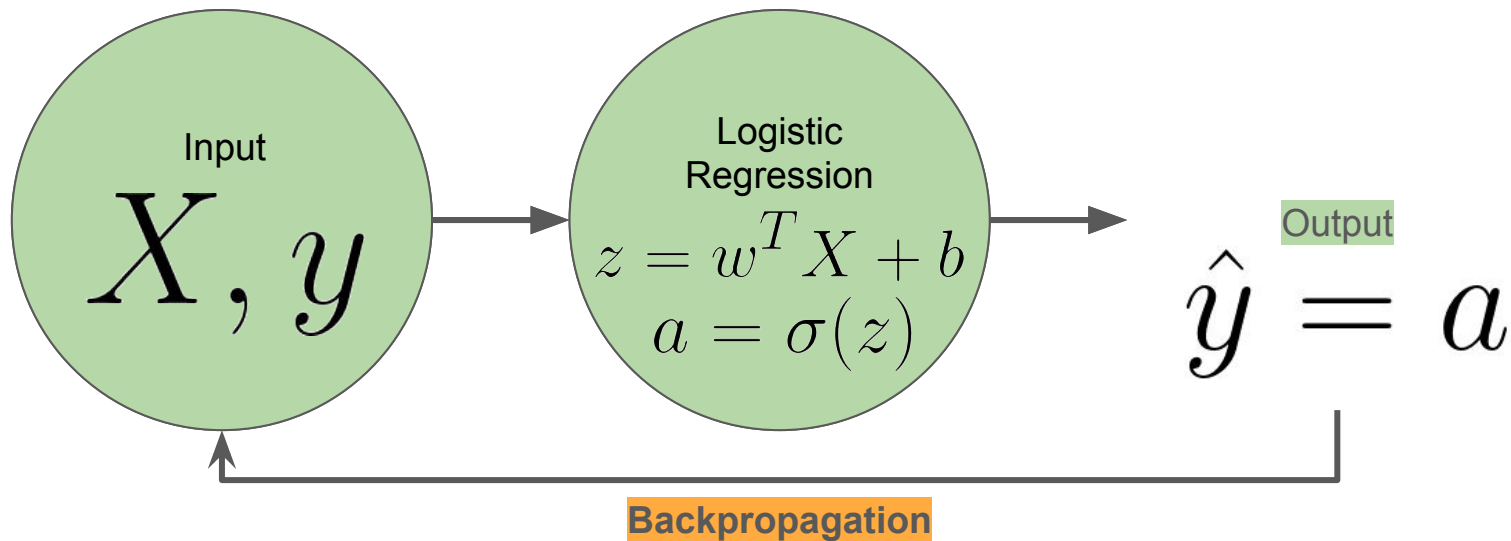
Questions?

But how do we get good predictions?



How to choose good w, b ?

Second step: backward propagation



Two parts to backprop:

- Calculating loss / cost
 - Cost function is higher when model performs poorly, lower when model performs well
- Gradient descent
 - Use derivatives to adjust the values of w and b

Cost of one example breaks down into two cases:

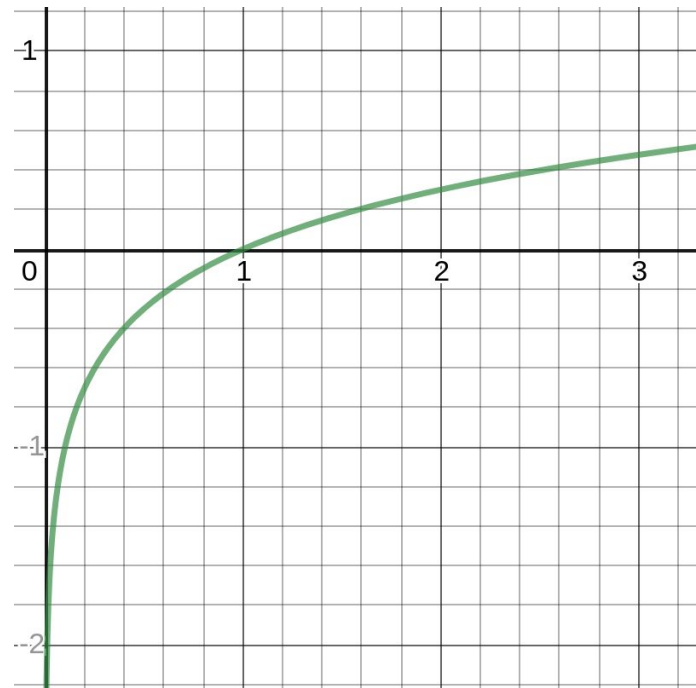
- If $y^{(i)} = 1$:
 - If $\hat{y}^{(i)} \approx 1$, model did well and should have low cost
 - If $\hat{y}^{(i)} \approx 0$, model did poorly and should have high cost
- If $y^{(i)} = 0$:
 - If $\hat{y}^{(i)} \approx 0$, model did well and should have low cost
 - If $\hat{y}^{(i)} \approx 1$, model did poorly and should have high cost

Cost of one example is a piecewise function

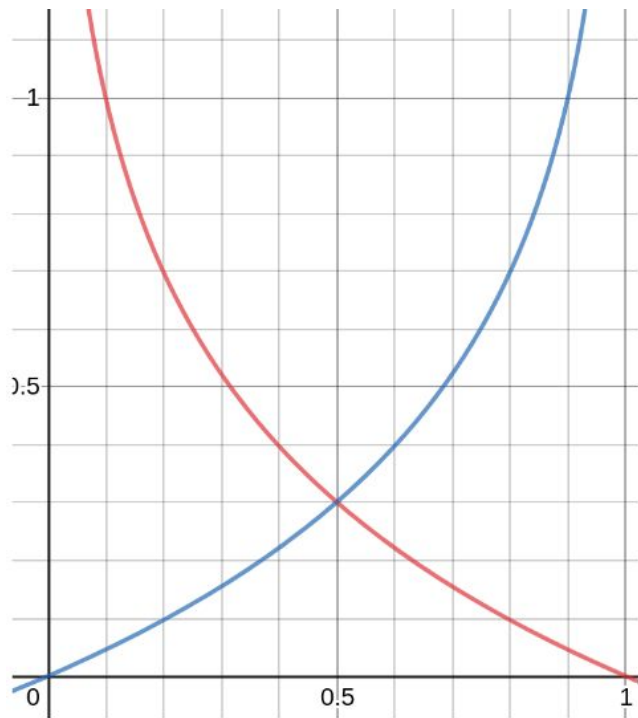
- We can cleverly use logarithms for our costs:

$$L(\hat{y}, y) = \begin{cases} y = 1 & -\log(\hat{y}) \\ y = 0 & -\log(1 - \hat{y}) \end{cases}$$

- What about $y_{\text{hat}}^{(i)} > 1$ or $y_{\text{hat}}^{(i)} < 0$?



Cost of one example looks like this when graphed



- Red: $y = 1$
- Blue: $y = 0$

$$L(\hat{y}, y) = \begin{cases} y = 1 & -\log(\hat{y}) \\ y = 0 & -\log(1 - \hat{y}) \end{cases}$$

Total cost should take into account all examples

- We simply take the average:

$$J(w, b) = \frac{1}{m} \sum_{i=1}^m L(\hat{y}^{(i)}, y^{(i)})$$

- J can be represented in terms of w and b

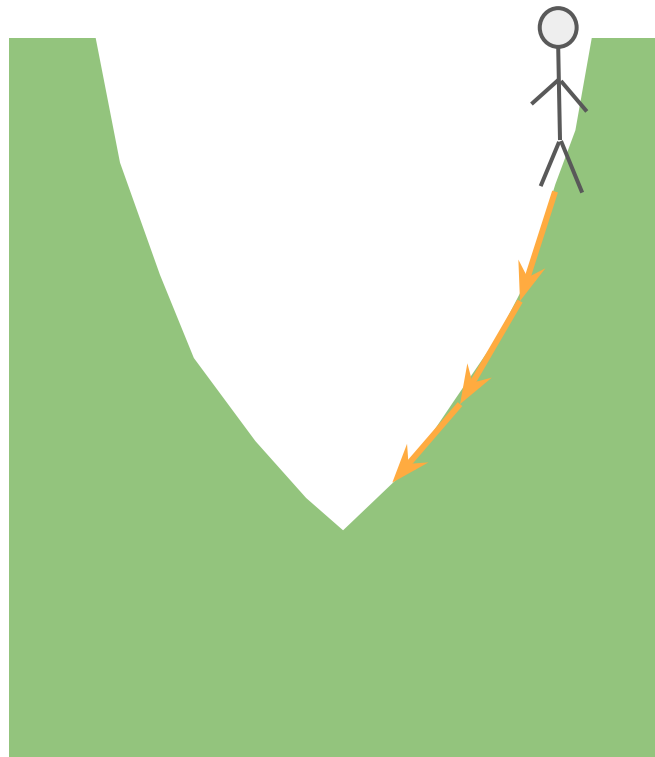
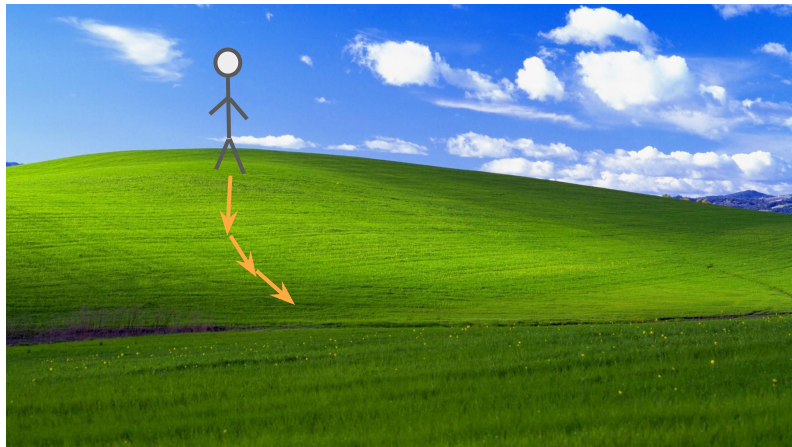
Note: Cost on one example is usually written like so

$$L(\hat{y}, y) = -(y \log(\hat{y}) + (1 - y) \log(1 - \hat{y}))$$

- This is effectively the same as the piecewise function, but is quicker to implement in code and makes computation faster
 - Plug in values for y to see why

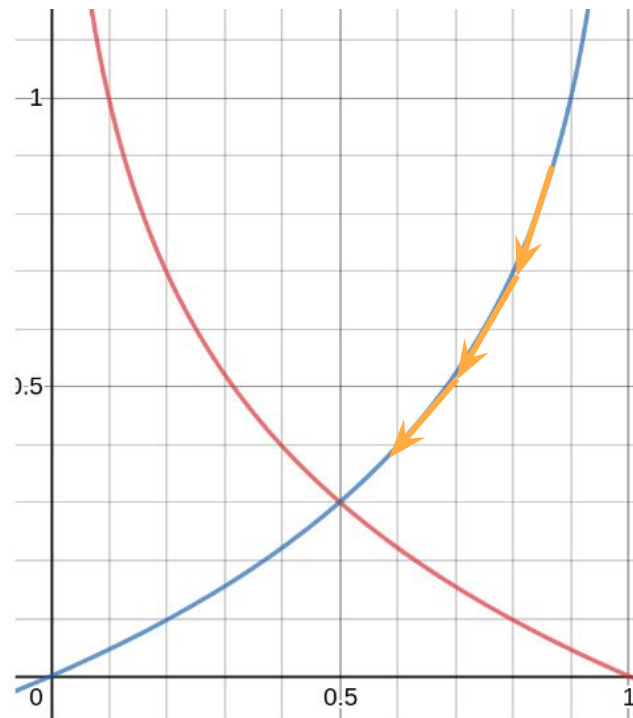
Gradient descent is like descending a hill

- Objective is to get to the bottom
- Go down in the direction of steepest slope
 - What would this look like in 3 dimensions?
 - What about > 3 ?



Gradient descent is like descending a hill

- We calculate slope using derivatives
 - “slope at a point, based on some variable”
- Subtract the current value by the derivative
- We will not prove the correctness of the derivatives – it’s more important to know why we’re using them



We calculate 3 important derivatives

- All derivatives are with respect to the total cost function J
 - J depends on multiple variables, which means we're really calculating partial derivatives
- As shorthand, we write only the denominators of the derivatives
- We use:
 - dz , an intermediate value
 - dw , used to update w
 - db , used to update b

dz

- Simply subtract labels y from predictions a

$$dz = a - y$$

$d\mathbf{w}$

- X is (n, m) while $d\mathbf{z}$ is $(1, m)$
 - Transposed to $(m, 1)$

$$d\mathbf{w} = \frac{1}{m} (X d\mathbf{z}^T)$$

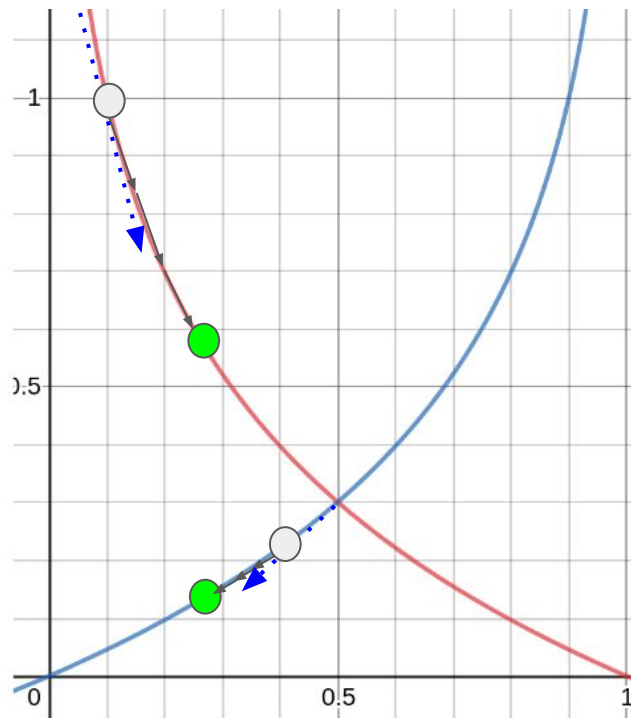
db

- Average of the elements in dz

$$db = \frac{1}{m} \sum_{i=1}^m dz^{(i)}$$

From variable values, subtract their derivatives

- Moving down the hill
 - The “bottom” of the hill is the minimum of the cost function
- We usually scale down the size of this step
 - Learning rate, α



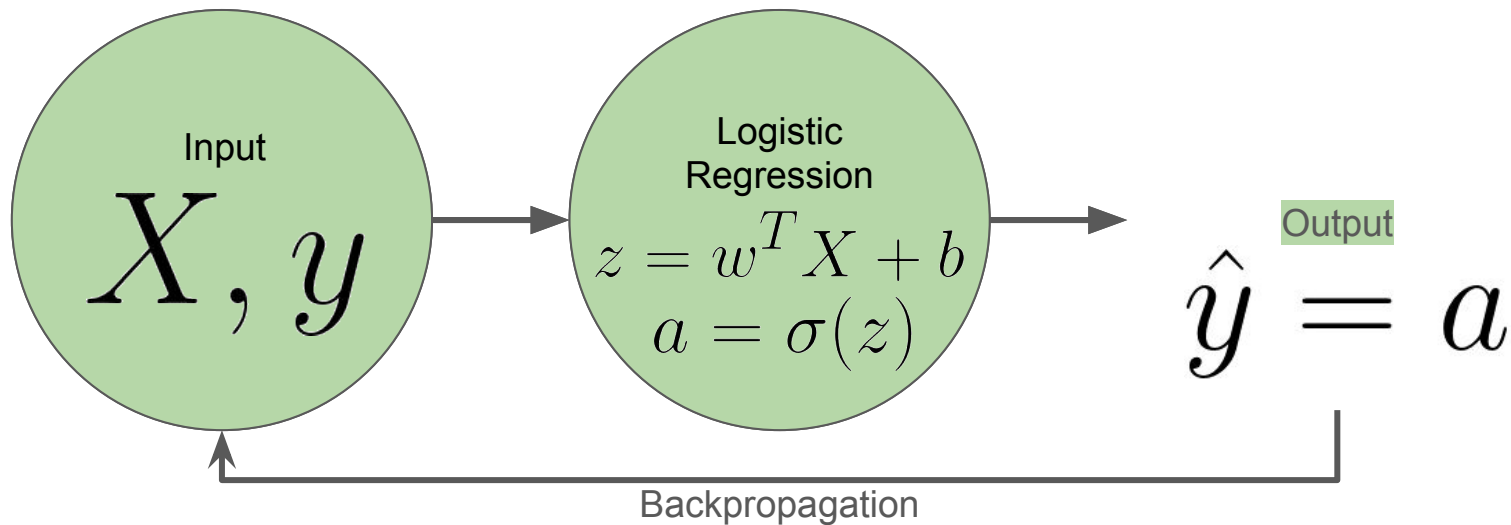
We update w and b with the derivatives

$$w := w - \alpha dw$$

$$b := b - \alpha db$$

- No reason to update z because it depends on w and b

And that's the entire process!



- That is one iteration – now do it for thousands of iterations

Review

- Input:

- Training set X with m instances $x^{(i)}$
 - Each instance has n features x_1, x_2, \dots
- Labels y with m labels $y^{(i)}$

- Output:

- Predicted probabilities \hat{y} with m predictions $\hat{y}^{(i)}$

Review

- Logistic regression:
 - Weight vector w and bias value b
 - Use these to calculate z
 - Input z into the activation function, the logistic / sigmoid function
 - The value we get is a
 - Fine-tune the weights by running gradient descent
 - Subtracting slopes (partial derivatives)

Table of important values

Name	Calculated during:	Dimensions	Purpose
X	Input	(n,m)	Input dataset
y	Input	$(1,m)$	Labels
$y_{\hat{}} (=a)$	Forward propagation	$(1,m)$	Predictions made by model
w	Backward propagation	$(n,1)$	Weights, used in forward prop
b	Backward propagation	$(1,1)$, scalar	bias, used in forward prop
z	Forward propagation	$(1,m)$	Intermediate value in forward prop
J	Backward propagation	$(1,1)$, scalar	Measures model inaccuracy

Questions?